

▼ Machine Learning Lab 4EII - IA course

1- Introduction

This lab aim at using simple **machine learning** classifiers via sklearn Python module for a classification problem of recognizing handwritten digits. We consider ten digits 0 to 9 from MNIST dataset, with using some machine learning algorithms, i.e. **Decision Trees, Random Forest** and **svm** used as **classifiers**.

In this lab you will learn to:

- Use the sklearn machine learning library
- Use different machine learning algorithms to classify 10 digits from MNIST dataset
- Study the performance and hyper-parameters of these classifiers
- Display graph results using the matplotlib module and confusion matrix

▼ 2- Module importation

Import some useful and common python modules

```
import numpy as np
import matplotlib.pyplot as plt
from sklearn.datasets import fetch_openml
import progressbar
import time
```

▼ 3- Download and study the the MNIST dataset

▼ 3.a - Download the MNIST dataset

MNIST dataset contains 70000 images of handwritten digits from 0 to 9. The dataset contains images of size 28x28 pixels and the corresponding labels

```
mnist = fetch_openml('mnist_784') #You can also use and test Fashion-MNIST dataset
```

▼ 3.b - Create a class structure to save and analyse the dataset

```
def computeentropy(image):
    lensig=image.size
    symset=list(set(image))
```

```

numsym=len(symset)
propab=[np.size(image[image==i])/(1.0*lensig) for i in symset]
ent=np.sum([p*np.log2(1.0/p) for p in propab])
return ent;

class Digit:
    def __init__(self, data, target):
        self.width = int(np.sqrt((len(data))))
        self.target = target;
        self.image = data;
        self.features = {
            'var'          :0.0, 'std'          :0.0,
            'mean'         :0.0, 'entropy'      :0.0,
        }
        self.computeFeatures()

    def computeFeatures(self):
        self.features['var'] = round(np.var(self.image),2)
        self.features['std'] = round(np.std(self.image),2)
        self.features['mean'] = round(np.mean(self.image),2)
        self.features['entropy'] = round(computeentropy(self.image),2)

    def print(self):
        print("Digit target: " + str(self.target))
        print("Digit target size: "+ str(self.width) + "x" +str(self.width) +
            '| mean : ' + str(self.features['mean']) +
            '| var : ' + str(self.features['var']) +
            '| std : ' + str(self.features['std']) +
            '| entropy : ' + str(self.features['entropy']))
        print("Digit image:")
        plt.figure()
        plt.gray()
        plt.matshow(self.image.reshape(self.width, self.width))
        plt.savefig(str(self.target)+'.png', bbox_inches='tight')
        plt.show()

class Dataset:
    def __init__(self, data, size=0):
        self.length = int((len(data['data'])))
        if size > 0 and size < self.length:
            self.length = size;
        else:
            size = self.length;

        self.targets = data['target'][0:size]
        self.data = data['data'][0:size];
        self.digits = []
        self.createDigits()
        self.X_train = [];
        self.X_test = [];
        self.y_train = [];
        self.y_test = [];

    def printInfo(self):
        from collections import Counter

```

```

c = Counter(self.targets)
info = "Dataset size " + str(self.length)
key_value = {}
for i in sorted(c.keys()):
    key_value[i] = c[i];

plt.bar(key_value.keys(), key_value.values());
plt.xlabel('Labels')
plt.ylabel('Occurrence')
plt.title('Occurrence of MNIST dataset labels')
ax = plt.axes()
ax.grid(which='major', axis='y')
plt.show()
return info

def createDigits(self):
    for i in range(self.length):
        self.digits.append(Digit(self.data[i], self.targets[i]))

def separate_train_test(self, test_size_ratio):
    from sklearn.model_selection import train_test_split

    self.X_train, self.X_test, self.y_train, self.y_test = train_test_split(self.data, self.targets, test_size=test_size_ratio)
    # data normalization
    self.X_train = self.X_train/255;
    self.X_test = self.X_test/255;
    print('Size of training set : ' + str(len(self.y_train)) + ' / ' + str(len(self.digits)))
    print('Size of testing set : ' + str(len(self.y_test)) + ' / ' + str(len(self.digits)))

def display_train_test(self):
    from collections import Counter

    test = Counter(self.y_test)
    train = Counter(self.y_train)
    info = "Dataset size " + str(self.length)

    key_value_train = {};
    key_value_test = {};

    for i in sorted(test.keys()):
        key_value_test[i] = test[i];
    for i in sorted(train.keys()):
        key_value_train[i] = train[i];

    p1 = plt.bar(key_value_train.keys(), key_value_train.values(), width=0.5);
    p2 = plt.bar( key_value_test.keys(), key_value_test.values(), width=0.5, bottom=1);

    plt.legend((p1[0], p2[0]), ('Training set', 'Test set'), loc='lower left')
    plt.xlabel('Labels')
    plt.ylabel('Occurrence')
    plt.title('Occurrence of training and testing sets')
    ax = plt.axes()
    ax.grid(which='major', axis='y')
    plt.show();

```

▼ 3.b - Load the MNIST dataset in Dataset class and analyse it:

1. Load the dataset in Dataset class

samples is the number of considered samples (sub-set) over 700000 of MNIST dataset, it enables faster training and testing

```
samples = 20000;  
#TO BE COMPLETED  
training_set = Dataset(mnist, samples)
```

2. Display some digist with corresponding features


```
samples_to_diply = 10#TO BE COMPLETED  
for i in range(samples_to_diply):  
    training_set.digits[i].print()
```

Digit target: 5

Digit target size: 28x28| mean : 35.11| var : 6343.94| std :79.65| entropy :1

Digit image:

<Figure size 432x288 with 0 Axes>




Digit target: 0

Digit target size: 28x28| mean : 39.66| var : 7037.06| std :83.89| entropy :1

Digit image:

<Figure size 432x288 with 0 Axes>




Digit target: 4

Digit target size: 28x28| mean : 24.8| var : 4300.7| std :65.58| entropy :1.4

Digit image:

<Figure size 432x288 with 0 Axes>




Digit target: 1

Digit target size: 28x28| mean : 21.86| var : 4366.42| std :66.08| entropy :1

Digit image:

<Figure size 432x288 with 0 Axes>

<Figure size 432x288 with 0 Axes>




Digit target: 9

Digit target size: 28x28| mean : 29.61| var : 5531.09| std :74.37| entropy :1

Digit image:

<Figure size 432x288 with 0 Axes>




Digit target: 2

Digit target size: 28x28| mean : 37.76| var : 6577.97| std :81.1| entropy :2.

Digit image:

<Figure size 432x288 with 0 Axes>





Digit target: 1

Digit target size: 28x28| mean : 22.51| var : 4602.49| std :67.84| entropy :0

Digit image:

<Figure size 432x288 with 0 Axes>






Digit target: 3

Digit target size: 28x28| mean : 45.75| var : 8102.99| std :90.02| entropy :1

Digit image:

<Figure size 432x288 with 0 Axes>



Digit target: 1

Digit target size: 28x28| mean : 13.87| var : 2768.36| std :52.62| entropy :0

3. Display digits repartitions with *printInfo* function of *Dataset* class

- Is the dataset well balanced ?




As we can see, there is a smooth discrepancy related to the label '1', but nothing relevant enough to invalid the dataset.



```
training_set.printInfo()
```

/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:70: MatplotlibDe



▼ 4 - Dataset preparation

The MNIST dataset is split to training and testing sets with the corresponding labels

0 1 2 3 4 5 6 7 8 9

▼ 4.a - Split the the MNIST dataset in training and testing sets

- Use *separate_train_test* function with a test set split ratio as parameter
- The test and train sets will be loaded in X_train and X_test lists and the corresponding labels in y_train and y_test lists.

```
test_ratio = 0.2;
#TO BE COMPLETED
training_set.separate_train_test(test_ratio)
```

Size of training set : 16000 / 20000
Size of testing set : 4000 / 20000

▼ 4.b - Display the repartition of the digits

- Use *display_train_test* function to illustrate the digits' repartition
- Check whether the repartition ratio is correct

```
#TO BE COMPLETED
training_set.display_train_test()
```


/usr/local/lib/python3.7/dist-packages/ipykernel_launcher.py:111: MatplotlibD

Occurrence of training and testing sets

Since we have 20000 samples for label '0' and the test ratio is 25%, we should have 4000 test samples selected, which can be noticed through the graph above. The proportion is maintained for the rest of the dataset as well.



▼ 5 - Classifier Training and testing



▼ 5.a - Training and testing the Decision Tree (DT) model

In this section, you will have to initialize a Decision Tree classifier and train it with the generated training set.

Decision Trees (DTs) are a non-parametric supervised learning method used for classification and regression. The goal is to create a model that predicts the value of a target variable by learning simple decision rules inferred from the data features.

Useful functions \Rightarrow `DecisionTreeClassifier` and `DecisionTreeClassifier.fit()`

For more details, you can refer to the sklearn documentation [Decision Tree](#)

```
from sklearn.tree import DecisionTreeClassifier
depth = 5;
clf = DecisionTreeClassifier(random_state=0, max_depth=depth)

clf.fit(training_set.X_train, training_set.y_train)
print ("Accuracy on training set " + str(round(clf.score(training_set.X_train, training_set.y_train), 2)))
print ("Accuracy on testing set " + str(round(clf.score(training_set.X_test, training_set.y_test), 2)))

Accuracy on training set 0.69
Accuracy on testing set 0.68
```

▼ 5.b - Hyper-parameters optimisation of the Decision Tree (DT) model

In this section you will train the DT model with different depths and select the one that enables the best performance in terms of trade-off between accuracy on the testing test and complexity while avoiding over-fitting

```
from sklearn.tree import DecisionTreeClassifier

depths = [2, 5, 8, 10, 15, 20, 25]
scores_training = []
for depth in depths:
    clf = DecisionTreeClassifier(max_depth=depth, random_state=0)
    clf.fit(training_set.X_train, training_set.y_train)
    scores_training.append(clf.score(training_set.X_train, training_set.y_train))
```

```

score_training = [0.0 for i in range(len(depths))]
score_testing  = [0.0 for i in range(len(depths))]
time_train     = [0.0 for i in range(len(depths))]
time_test      = [0.0 for i in range(len(depths))]
idx=0;
bar = progressbar.ProgressBar(maxval=len(depths)).start()

for d in depths:
    clf = DecisionTreeClassifier(random_state=0, max_depth=d)
    t = time.process_time()
    #TO BE COMPLETED perform training
    clf.fit(training_set.X_train, training_set.y_train)
    time_train[idx] = time.process_time() - t
    score_training[idx] = round(clf.score(training_set.X_train, training_set.y_train),2)
    score_testing[idx] = round(clf.score(training_set.X_test, training_set.y_test),2)
    time_test[idx] = time.process_time() - t - time_train[idx]
    idx +=1;
    bar.update(idx)

plt.figure(figsize=(8,4))
fig, ax1 = plt.subplots()


color = 'tab:blue'
ax1.set_xlabel('Depth', fontsize=15,)
ax1.set_ylabel('Accuracy', color=color)
ax1.plot(depths, score_training, '-bo', label='Training set')
ax1.plot(depths, score_testing, '--b*', label='Testing set')
ax1.tick_params(axis='y', labelcolor=color)
plt.grid()
plt.legend(loc='upper left')
ax2 = ax1.twinx() # instantiate a second axes that shares the same x-axis
color = 'tab:red'

ax2.set_ylabel('time (s)', color=color)
ax2.plot(depths, time_train, '-r<', label='Training set') # we already handled the
ax2.plot(depths, time_test, '--r>', label='Testing set') # we already handled the
ax2.tick_params(axis='y', labelcolor=color)
fig.tight_layout() # otherwise the right y-label is slightly clipped

plt.show()
plt.savefig('perf.png')

```

100% (7 of 7) |#####| Elapsed Time: 0:00:20 ETA: 00:00:
 <Figure size 576x288 with 0 Axes>

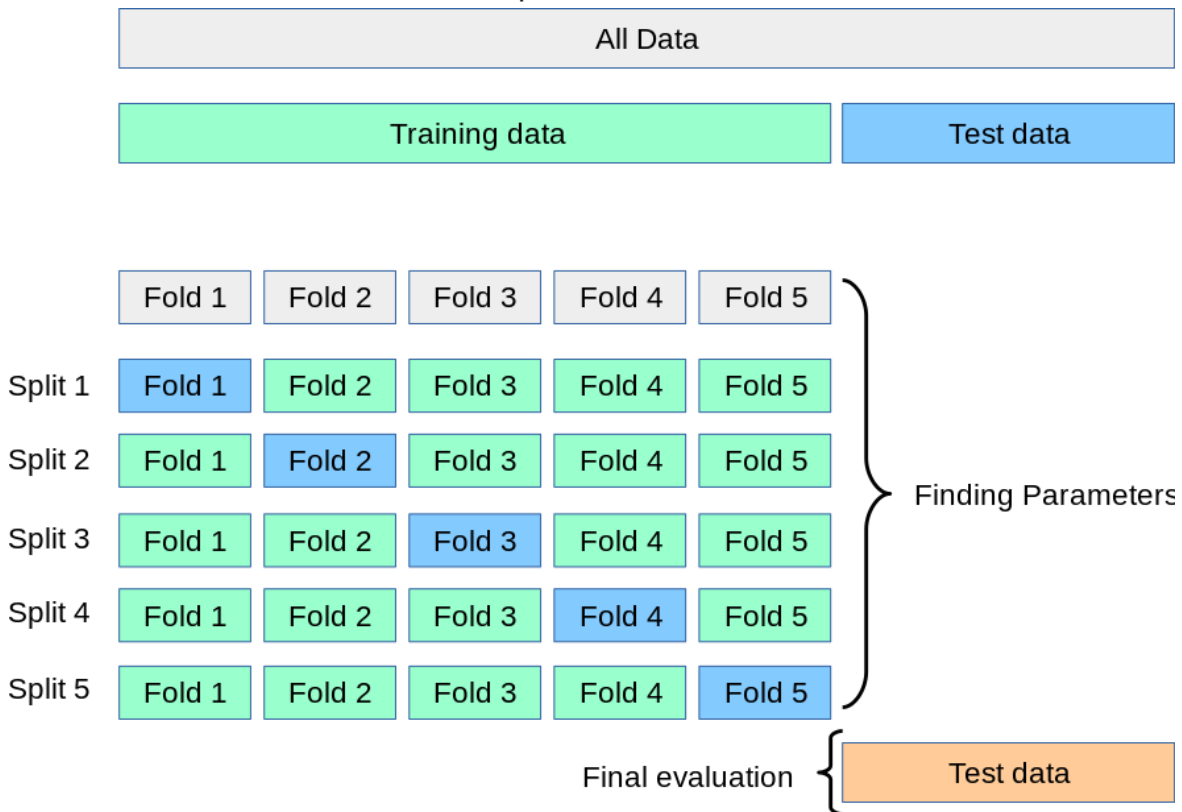


▼ 5.c - Cross-validation of the best performing solution

In this section you will test the best performing DT configuration in cross-validation approach.

The [cross-validation](#) will split the training set in k non-overlapping sub-sets. Then the model is trained on $(k-1)$ sub-sets and tested on the remaining sub-set. This process is performed k times on k different testing sub-sets and take the average accuracy with confidence interval (CI).

Illustration of the cross validation repartition.



Is the performance on testing set accurate and valid ?

```
from sklearn.model_selection import cross_val_score
best_depth = 15#TO BE COMPLETED
k = 3#TO BE COMPLETED Number of sub-sets < 4
clf = DecisionTreeClassifier(random_state=0, max_depth=best_depth)
clf.fit(training_set.X_train, training_set.y_train)
print ('Accuracy on training set= ' + str(round(clf.score(training_set.X_train, t
scores = cross_val_score(clf, training_set.X_train, training_set.y_train, cv=k, sc
print (scores)
print("Accuracy: %0.2f (CI : +/- %0.2f)" % (scores.mean(), scores.std() * 2))
```

```
print ('Accuracy on test set= ' + str(round(clf.score(training_set.X_test, train1
```

```
↳ Accuracy on training set= 0.99
   [0.82133483 0.79823739 0.81511344]
   Accuracy: 0.81 (CI : +/- 0.02)
   Accuracy on test set= 0.83
```

Considering the simplicity of the classifier, we can say that a score of 0.83% is accurate. Although, we expect higher accuracy on a reliable classifier.

▼ 5.d - Display the confusion matrix

By definition a confusion matrix C is such that C_{ij} is equal to the number of observations known to be in group i and predicted to be in group j .

1. `plot_confusion_matrix` function enables to display a confusion matrix cm

```
def plot_confusion_matrix(cm,
                          target_names,
                          title='Confusion matrix',
                          cmap=None,
                          normalize=True):
    import matplotlib.pyplot as plt
    import numpy as np
    import itertools

    accuracy = np.trace(cm) / float(np.sum(cm))
    misclass = 1 - accuracy

    if cmap is None:
        cmap = plt.get_cmap('Blues')
    plt.figure(figsize=(8, 6))
    plt.imshow(cm, interpolation='nearest', cmap=cmap)
    plt.title(title)
    plt.colorbar()

    if target_names is not None:
        tick_marks = np.arange(len(target_names))
        plt.xticks(tick_marks, target_names, rotation=45)
        plt.yticks(tick_marks, target_names)

    if normalize:
        cm = cm.astype('float') / cm.sum(axis=1)[:, np.newaxis]

    thresh = cm.max() / 1.5 if normalize else cm.max() / 2
    for i, j in itertools.product(range(cm.shape[0]), range(cm.shape[1])):
        if normalize:
            plt.text(j, i, "{:0.3f}".format(cm[i, j]),
                     horizontalalignment="center",
                     color="white" if cm[i, j] > thresh else "black")
        else:
            plt.text(j, i, "{:d}".format(cm[i, j]),
```

```

plt.text(j, i, '{:0.4f}'.format(cm[i, j]),
         horizontalalignment="center",
         color="white" if cm[i, j] > thresh else "black")
plt.tight_layout()
plt.ylabel('True label')
plt.xlabel('Predicted label\naccuracy={:0.4f}; misclass={:0.4f}'.format(accuracy
plt.show()

```

2. Compute the [confusion matrix](#) of the selected best performing solution and display it with `plot_confusion_matrix` function

which are the most difficult digits to predict ? Support your answer with numbers from the confusion matrix.

```



# Plot non-normalized confusion matrix
from sklearn.metrics import confusion_matrix
class_names = ['0', '1', '2', '3', '4', '5', '6', '7', '8', '9'];
titles_options = [("Confusion matrix, without normalization", False),
                  ("Normalized confusion matrix", True)]
y_hat = clf.predict(training_set.X_test);

#TO BE COMPLETED use confusion matrix function to create conf_mx
conf_mx = confusion_matrix(training_set.y_test, y_hat)

plt.figure(figsize=(10,6))
for title, normalize in titles_options:
    disp = plot_confusion_matrix(cm=conf_mx,
                                target_names=class_names,
                                title=title,
                                cmap=plt.cm.Blues,
                                normalize=normalize)

```

<Figure size 720x432 with 0 Axes>



The most difficult digit to label for this classifier was '8', since we have the lowest percentage (68.2%) of correct classifications. In contradiction, the digit '1' was the easiest to classify.

5 | 0.014 0.017 0.011 0.043 0.023 0.734 0.034 0.008 0.023 0.028 |

6 - Testing other ML models

This section you will test the performance (section 5) of other ML classifiers such as random forest and SVM.

| 0.014 0.017 0.011 0.043 0.023 0.734 0.034 0.008 0.023 0.028 |

6.a Random Forest classifier

A random forest is a meta estimator that fits a number of decision tree classifiers on various sub-samples of the dataset and uses averaging to improve the predictive accuracy and control over-fitting.

Initialize the model with `RandomForestClassifier(n_estimators=50,max_depth=7,random_state=0)`

The hyper-parameters are the number of trees (`n_estimators`) and maximum depth (`max_depth`).

For more details on Random Forest classifier in sklearn, you can refer to [Random Forest classifier](#)

```
from sklearn.ensemble import RandomForestClassifier

depths = [2,5,8,10,15,20]
trees = [1, 10, 20, 30, 40]

for s in trees:
    idx = 0;
    score_training = [0.0 for i in range(len(depths))]
    score_testing = [0.0 for i in range(len(depths))]
    time_train = [0.0 for i in range(len(depths))]
    time_test = [0.0 for i in range(len(depths))]
    print ("Number of trees of : " + str(s))
    bar = progressbar.ProgressBar(maxval=len(depths)).start()
    for d in depths:
        #TO BE COMPLETED initialize the RF model
        clf = RandomForestClassifier(n_estimators=s, max_depth=d)
        t = time.process_time()
        #TO BE COMPLETED run the training
        clf.fit(training_set.X_train, training_set.y_train)
        time_train[idx] = time.process_time() - t
        score_training[idx] = round(clf.score(training_set.X_train, training_set.y_train))
        time_test[idx] = time.process_time() - t - time_train[idx]
        score_testing[idx] = round(clf.score(training_set.X_test, training_set.y_test))
        idx +=1;
    bar.update(idx)

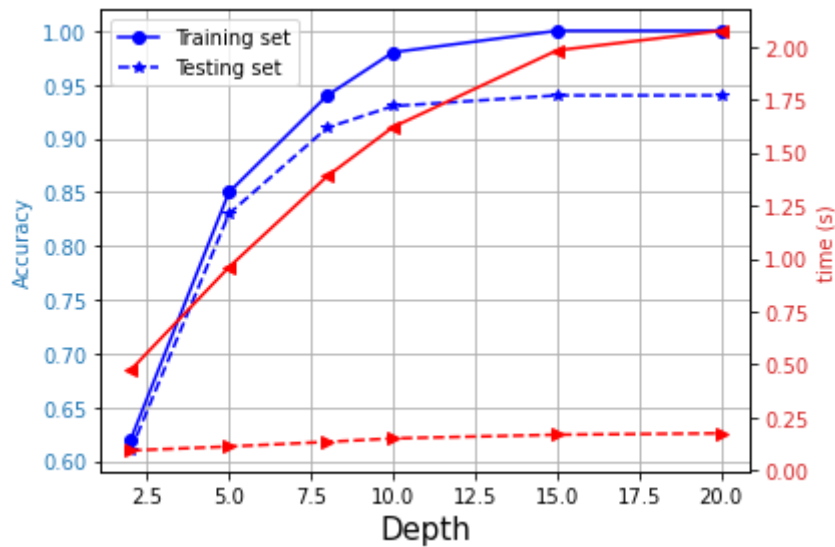
plt.figure(figsize=(8,4))
fig, ax1 = plt.subplots()

color = 'tab:blue'
ax1.set_xlabel('Depth', fontsize=15,)
ax1.set_ylabel('Accuracy', color=color)
ax1.plot(depths, score_training, '-bo', label='Training set')
ax1.plot(depths, score_testing, '--b*', label='Testing set')
ax1.tick_params(axis='y', labelcolor=color)
plt.legend(loc='upper left')
plt.grid()
ax2 = ax1.twinx() # instantiate a second axes that shares the same x-axis
color = 'tab:red'

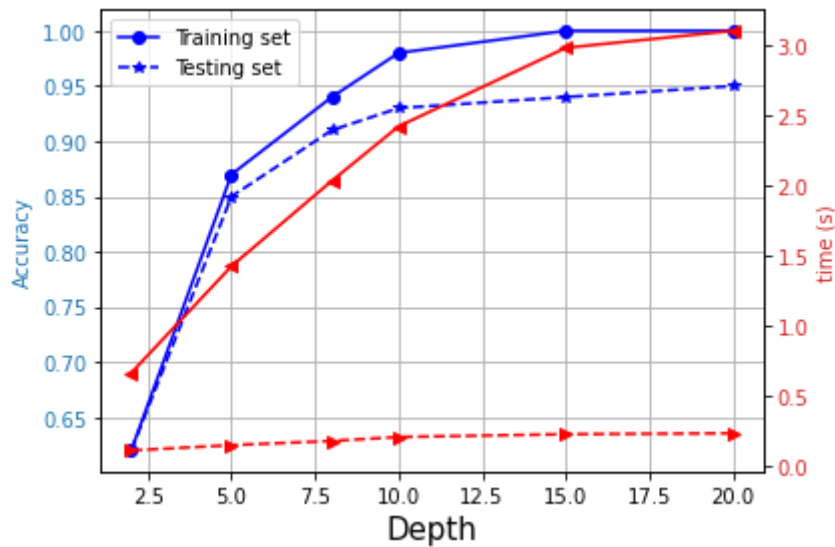
ax2.set_ylabel('time (s)', color=color)
ax2.plot(depths, time_train, '-r<', label='Training set') # we already handled
ax2.plot(depths, time_test, '--r>', label='Testing set') # we already handled
ax2.tick_params(axis='y', labelcolor=color)
fig.tight_layout() # otherwise the right y-label is slightly clipped

plt.show()
```

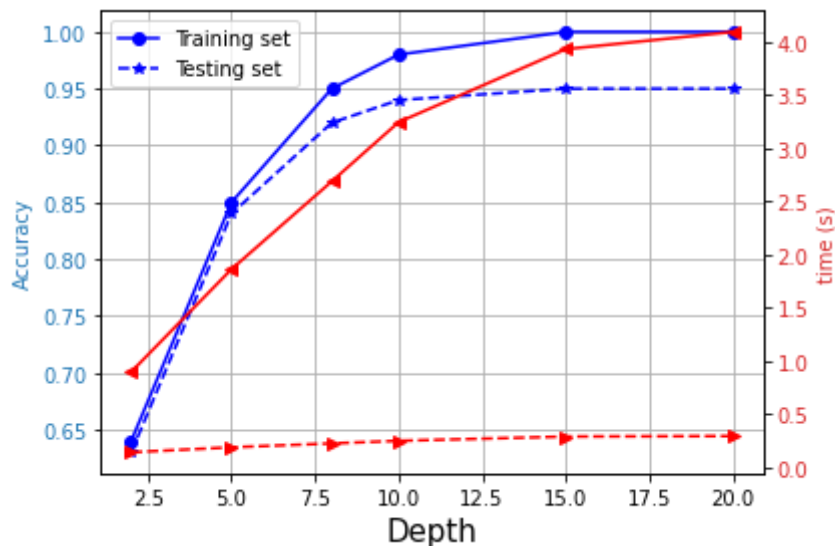
N/A% (0 of 6) | Elapsed Time: 0:00:00 ETA: --:--
 100% (6 of 6) | ##### Elapsed Time: 0:00:09 ETA: 00:00
 <Figure size 576x288 with 0 Axes>



N/A% (0 of 6) | Elapsed Time: 0:00:00 ETA: --:--
 100% (6 of 6) | ##### Elapsed Time: 0:00:14 ETA: 00:00
 <Figure size 576x288 with 0 Axes>



N/A% (0 of 6) | Elapsed Time: 0:00:00 ETA: --:--
 100% (6 of 6) | ##### Elapsed Time: 0:00:18 ETA: 00:00
 <Figure size 576x288 with 0 Axes>



▼ 6.b SVM classifier with different kernels

Initialize the SVM model by `svm.SVC(kernel=)` and test different kernels :

1. linear
2. poly (Polynomial) hyper-parameters: Optional degree=3-8 / gamma='scale'/auto
3. rbf (Radial Basis Function / Gaussian): Optional gamma='scale'/auto
4. sigmoid: Optional gamma='scale'/auto

For more details on svm classifier in sklearn, you can refer to [SVC classifier](#).

Compare your results with the performance reported on [MNIST](#) web site.

```
from sklearn import svm
kernels = ['linear', 'poly', 'rbf', 'sigmoid']
bar = progressbar.ProgressBar(maxval=len(kernels)).start()
idx = 0;
for ker in kernels:
    print('\nKernel: ' + ker)
```

```
#TO BE COMPLETED initialize the model
clf = svm.SVC(kernel=ker)
t = time.process_time()
#TO BE COMPLETED run the training
clf.fit(training_set.X_train, training_set.y_train)
print ( 'processing time : ' + str (round((time.process_time() - t),2)))
print ('Accuracy on training set= ' + str(round(clf.score(training_set.X_train,
print ('Accuracy on testing set= ' + str(round(clf.score(training_set.X_test, t
idx +=1;
bar.update(idx)
```

```
N/A% (0 of 4) | Elapsed Time: 0:00:00 ETA:  --:--:
Kernel: linear
processing time : 47.48
Accuracy on training set= 0.99
25% (1 of 4) |##### | Elapsed Time: 0:02:23 ETA: 0:07:

Kernel: poly
processing time : 78.92
Accuracy on training set= 0.99
50% (2 of 4) |##### | Elapsed Time: 0:05:27 ETA: 0:06:

Kernel: rbf
processing time : 75.61
Accuracy on training set= 0.99
75% (3 of 4) |##### | Elapsed Time: 0:09:01 ETA: 0:03:

Kernel: sigmoid
processing time : 84.64
Accuracy on training set= 0.8
100% (4 of 4) |#####| Elapsed Time: 0:13:10 ETA: 00:00:
```

We obtained better performance comparing 0.04 with the errors of 0.68, 0.68 e 0.56 for Virtual SVM deg-9 poly results.

7 - Test the performance of different classifiers with Fashion-MNIST dataset

You can find [here](#) the benchmark of different classifiers on Fashion-MNIST dataset

```
fashionMnist = fetch_openml('Fashion-MNIST') #You can also use and test Fashion-MNIST
```

```
samples = 15000
training_set = Dataset(fashionMnist, samples)
```

```
samples_to_display = 10#TO BE COMPLETED
for i in range(samples_to_display):
    training_set.digits[i].print()
```

